

Object detection with R-CNN family

Computer vision

Comparison between image classification, object detection and instance segmentation.

Classification




CAT

**Classification
+ Localization**



CAT

Object Detection



CAT, DOG, DUCK

**Instance
Segmentation**



CAT, DOG, DUCK

Single object

Multiple objects

Classification




Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.



Vector:
4096

→
Fully-Connected:
4096 to 1000

Class Scores


Cat: 0.9

Dog: 0.05

Car: 0.01

...

Classification + Localization



**Fully
Connected:**
4096 to 1000

Class Scores


Cat: 0.9
Dog: 0.05
Car: 0.01
...

Vector:
4096

**Fully
Connected:**
4096 to 4

**Box
Coordinates**
(x, y, w, h)

Classification + Localization



**Fully
Connected:**
4096 to 1000

Class Scores

Cat: 0.9
Dog: 0.05
Car: 0.01
...

Correct label:
Cat

**Softmax
Loss**


Vector:
4096
**Fully
Connected:**
4096 to 4

**Box
Coordinates**
(x, y, w, h)

L2 Loss

Correct box:
(x', y', w', h')

Classification + Localization



Fully Connected:
4096 to 1000

Class Scores

Cat: 0.9
Dog: 0.05
Car: 0.01
...

Correct label:
Cat

Softmax Loss

Multitask Loss

Vector:
4096

Fully Connected:
4096 to 4

Box Coordinates
(x, y, w, h)

L2 Loss

Correct box:
(x', y', w', h')

+ → **Loss**


Aside: Human pose estimation




Represent pose as a set of 14 joint positions:

- Left/right foot
- Left/right knee
- Left/right hip
- Left/right shoulder
- Left/right elbow
- Left/right hand
- Neck
- Head top


Object detection as regression?



CAT: (x, y, w, h)



DOG: (x, y, w, h)
DUCK: (x, y, w, h)




DOG: (x, y, w, h)
CAT: (x, y, w, h)
DUCK: (x, y, w, h)
DUCK: (x, y, w, h)

...


Object detection as regression?

Each image needs a different number of outputs!



CAT: (x, y, w, h)


4 numbers



DOG: (x, y, w, h)

8 numbers

DUCK: (x, y, w, h)



DOG: (x, y, w, h)

CAT: (x, y, w, h)

DUCK: (x, y, w, h)

DUCK: (x, y, w, h)

16+ numbers


...

Object detection as classification


Sliding window

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background.

Cat: NO
Dog: NO
Background: YES



Cat: YES
Dog: NO
Background: NO



Cat: NO
Dog: YES
Background: NO

Object detection as classification

Sliding window


Apply a CNN to many different crops of the image, CNN classifies each crop as object or background.

But is it good?

Cat: NO
Dog: NO
Background: YES



Cat: YES
Dog: NO
Background: NO



Cat: NO
Dog: YES
Background: NO


Object detection as classification

Sliding window

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background.

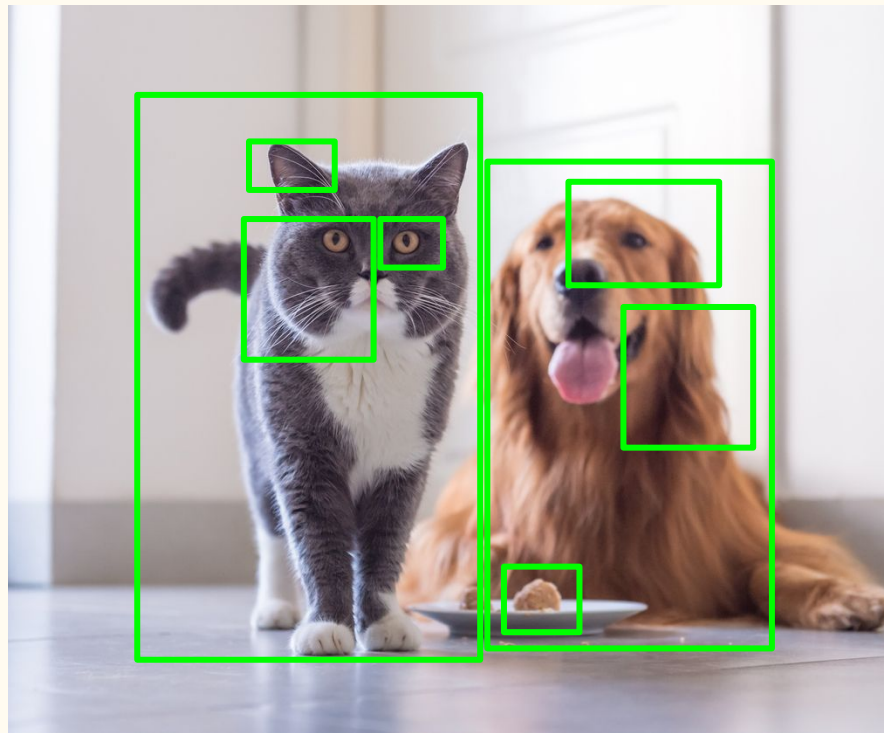
But is it good?

NO! Need to apply CNN to huge number of locations, scales, and aspect ratios, very computationally expensive



Region Proposals / Selective Search

- Find “blobby” image regions that are likely to contain objects
- Relatively fast to run;
e.g. Selective Search gives 2000 region proposals in a few seconds on CPU




R-CNN

R-CNN




Input image

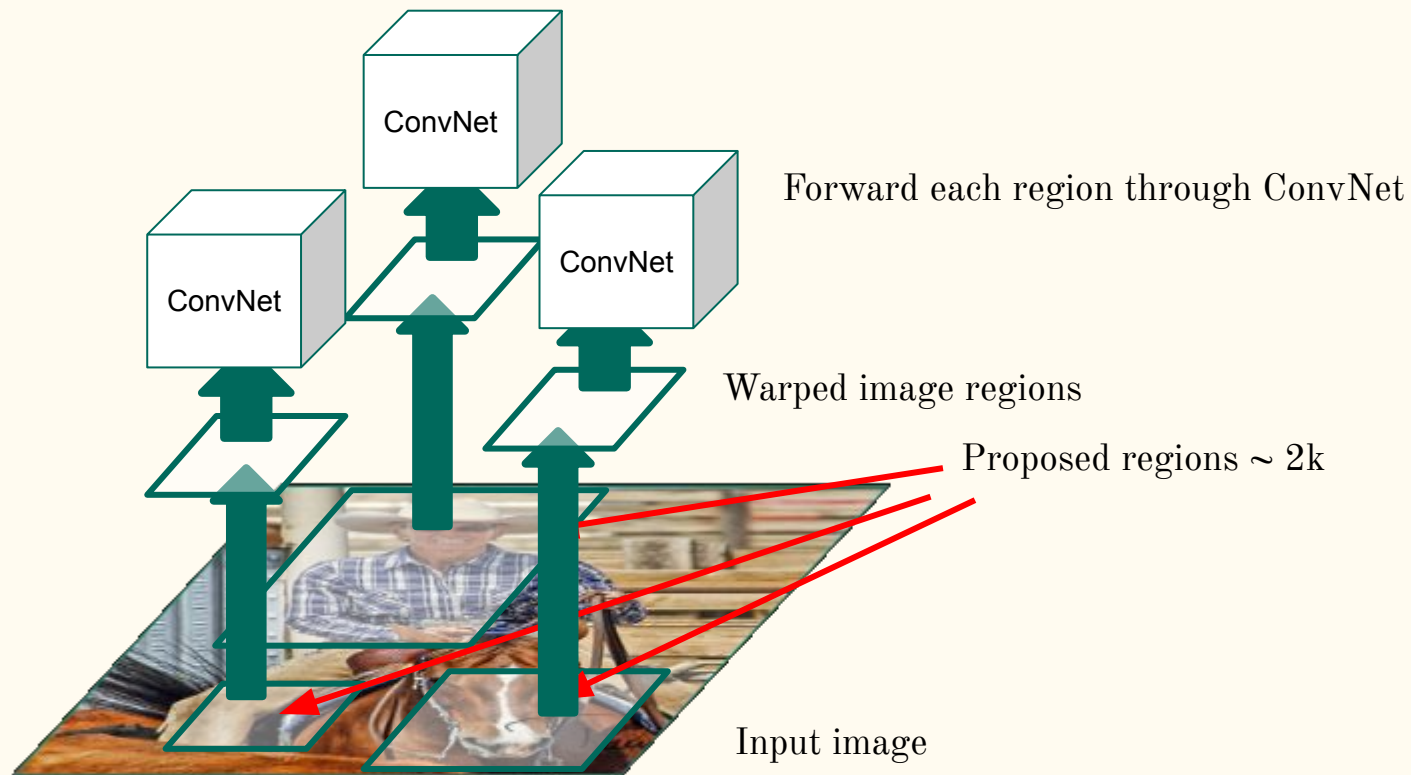
R-CNN



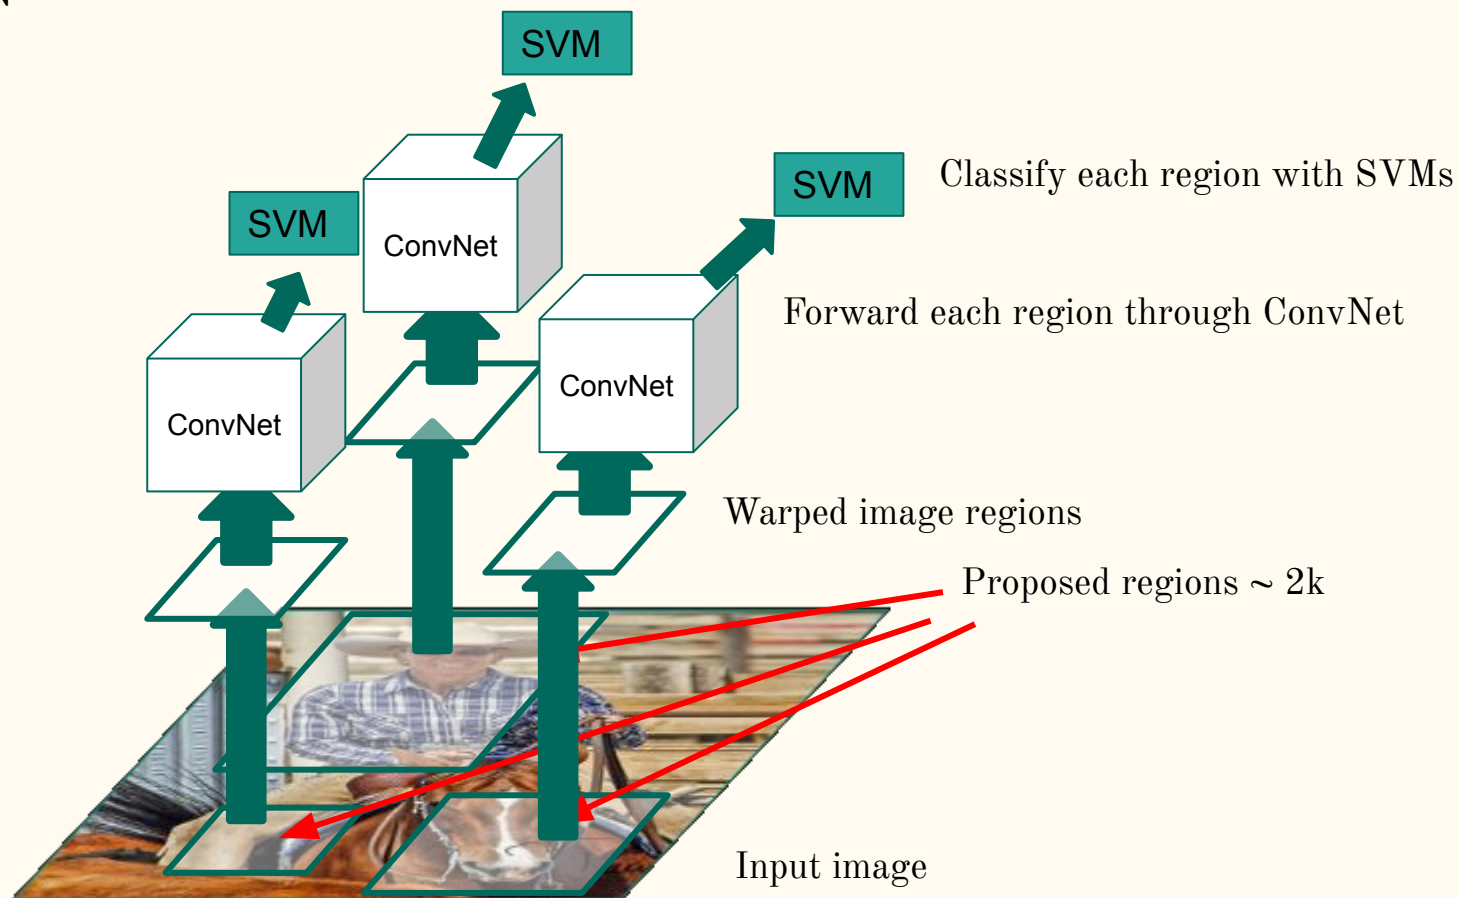
R-CNN



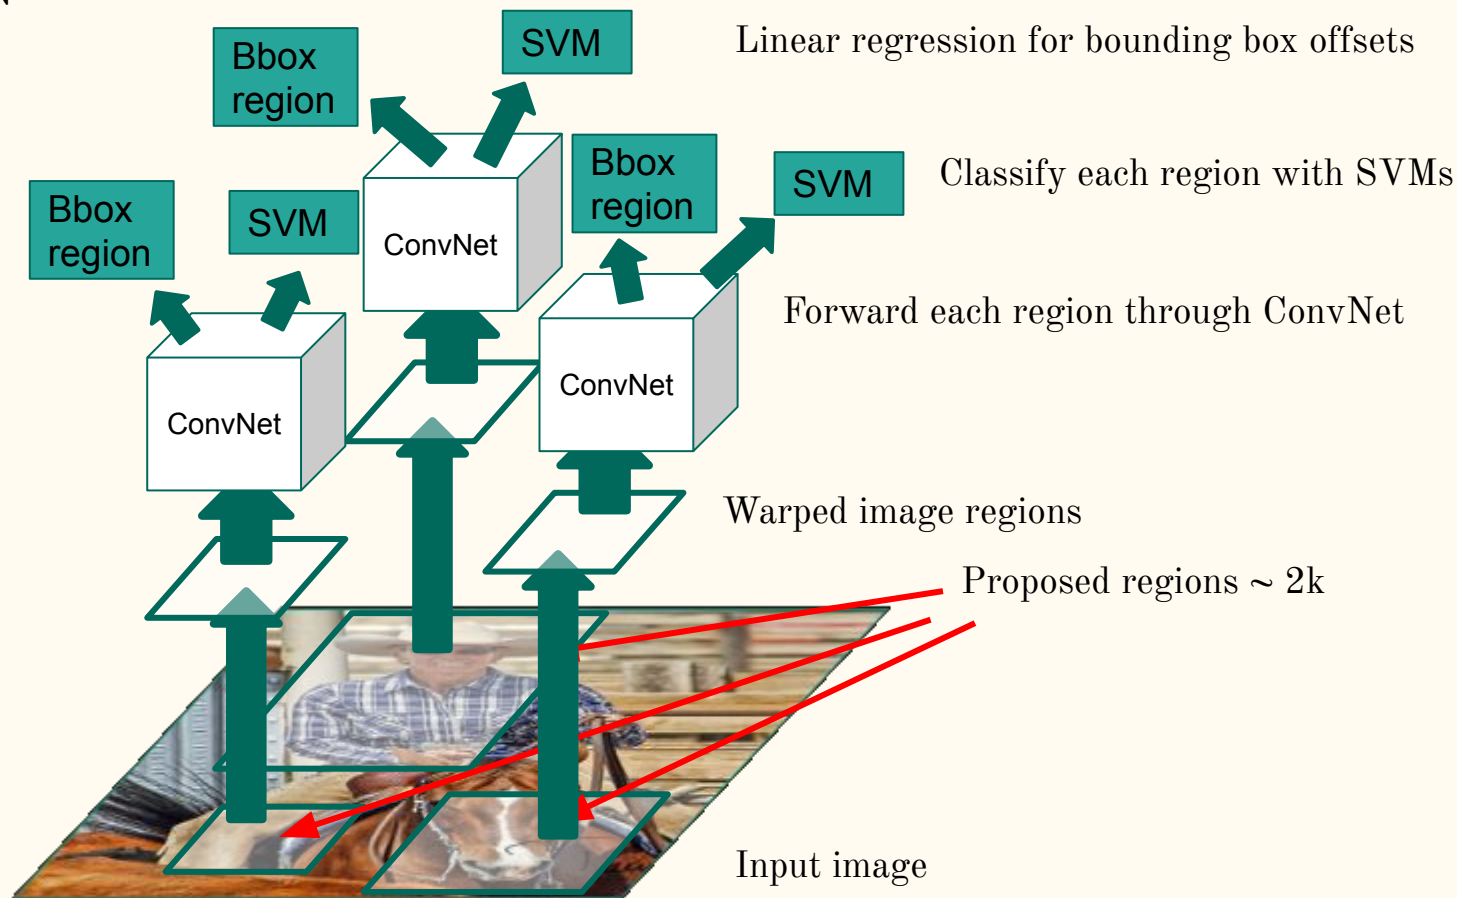
R-CNN



R-CNN




R-CNN



R-CNN problems

- Ad hoc training objectives
 - Fine-tune network with softmax classifier (log loss)
 - Train post-hoc linear SVMs (hinge loss)
 - Train post-hoc bounding-box regressions (least squares)
- Training is slow (84h), takes a lot of disk space
- Inference (detection) is slow
 - 47s / image with VGG16 [Simonyan & Zisserman. ICLR15]

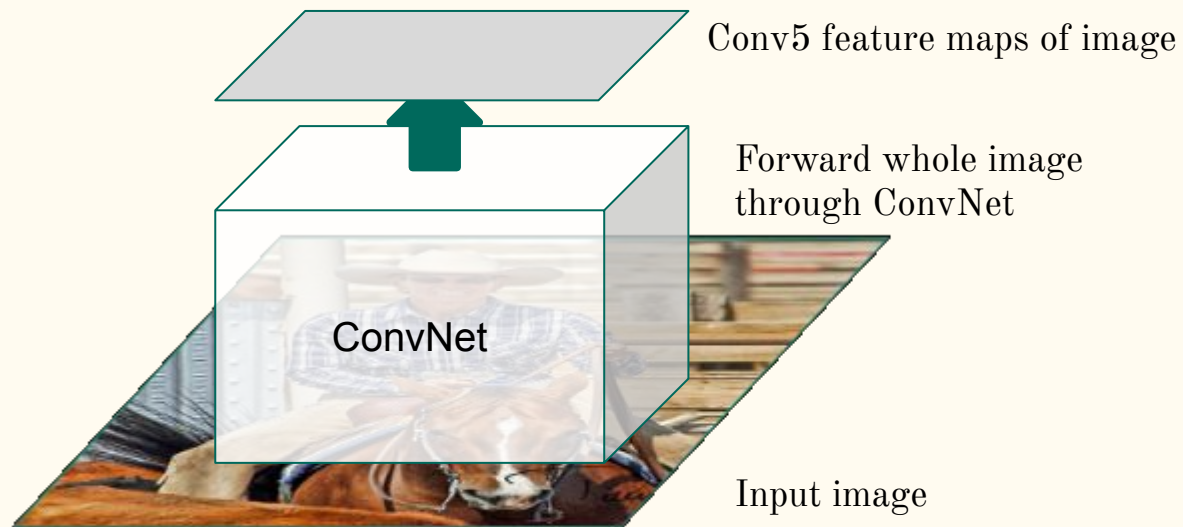


Fast R-CNN




Input image

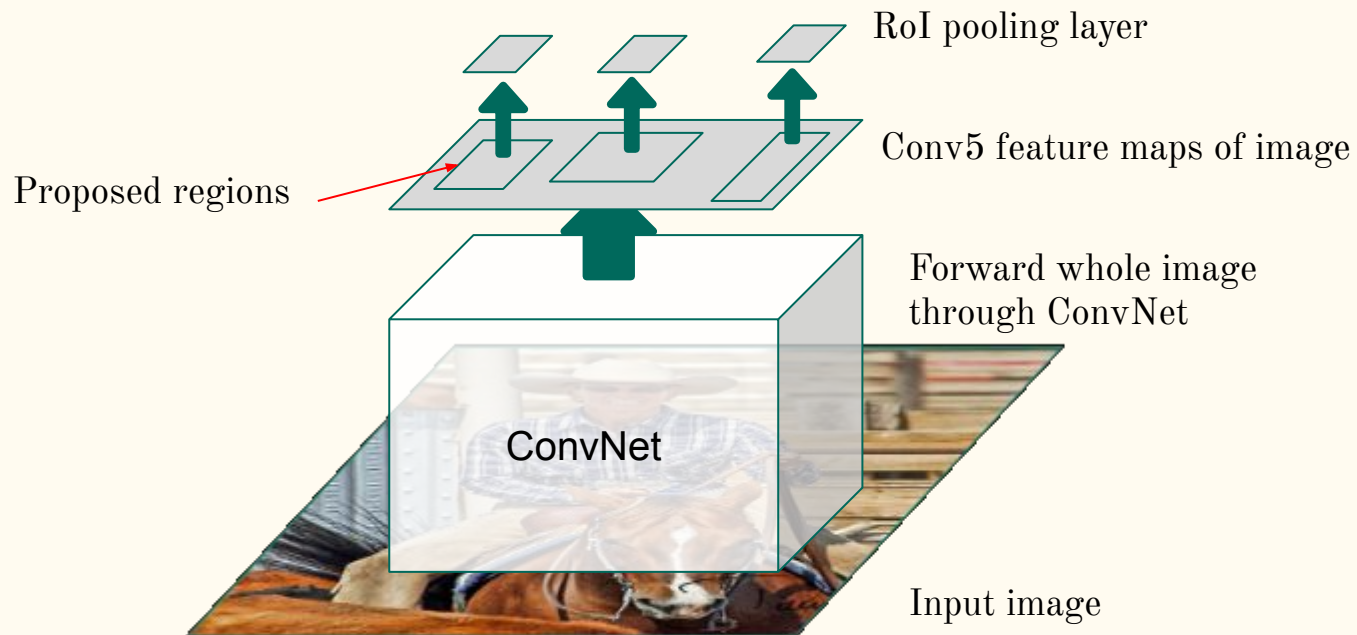
Fast R-CNN



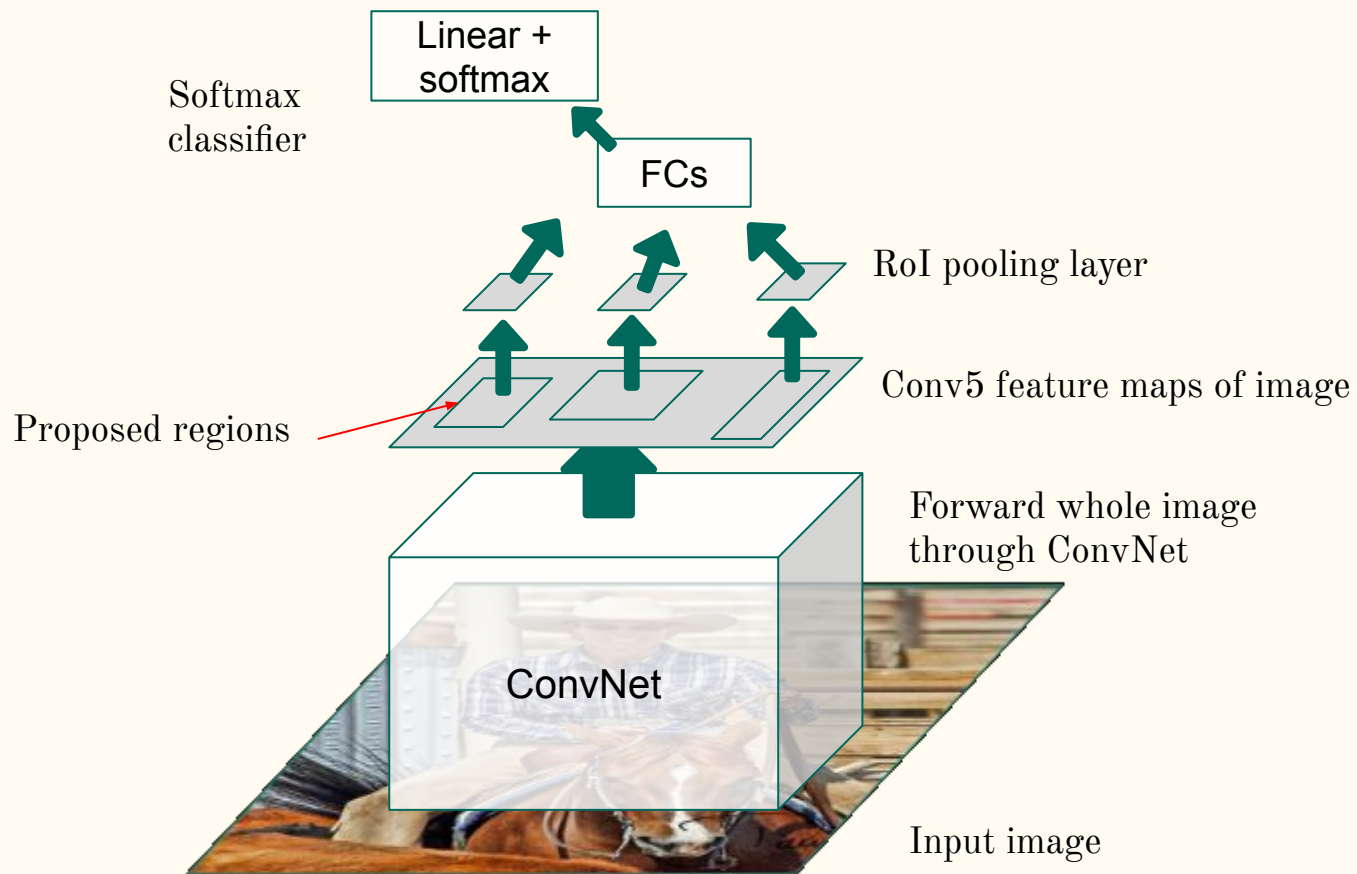
Fast R-CNN



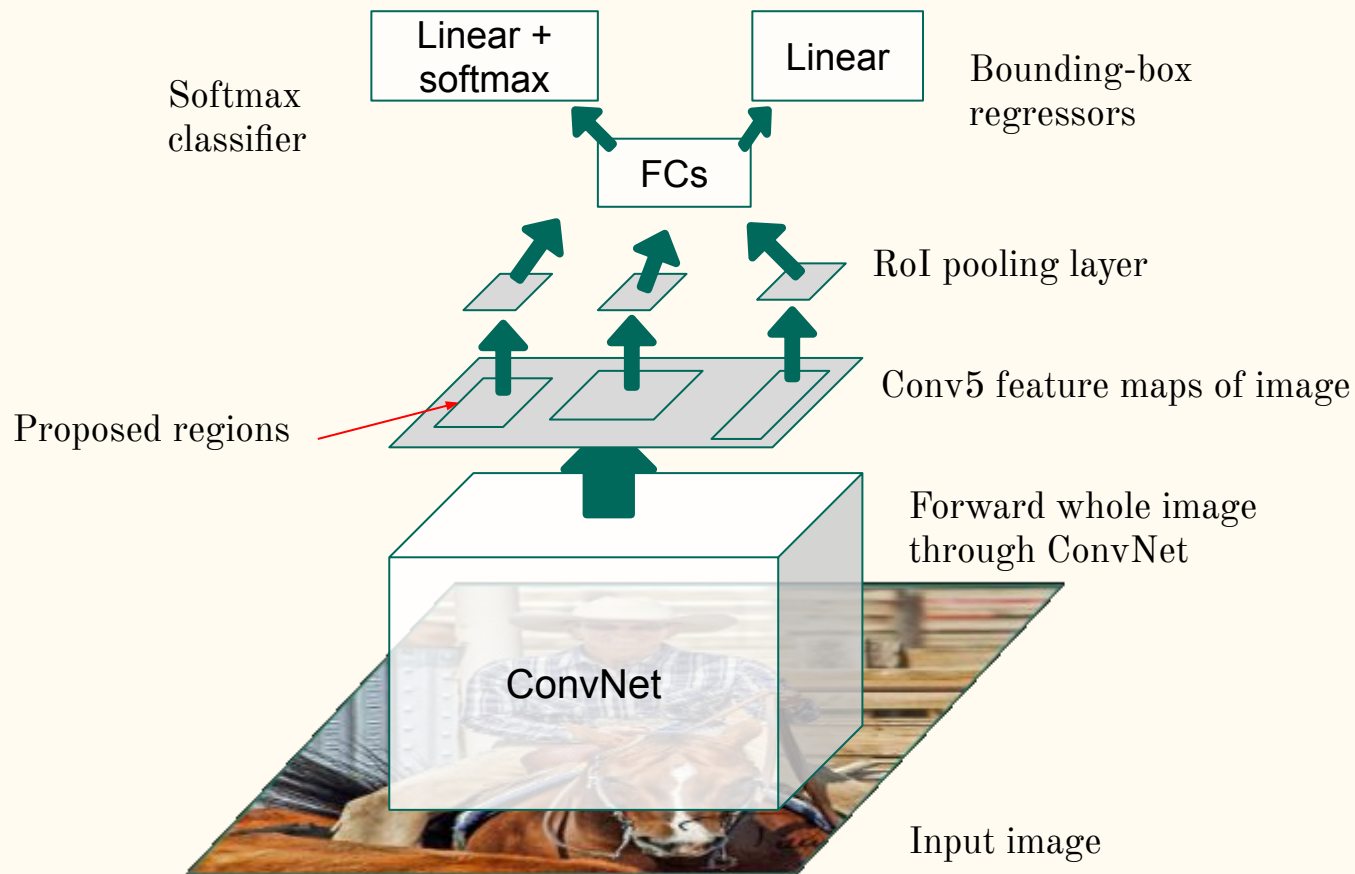
Fast R-CNN



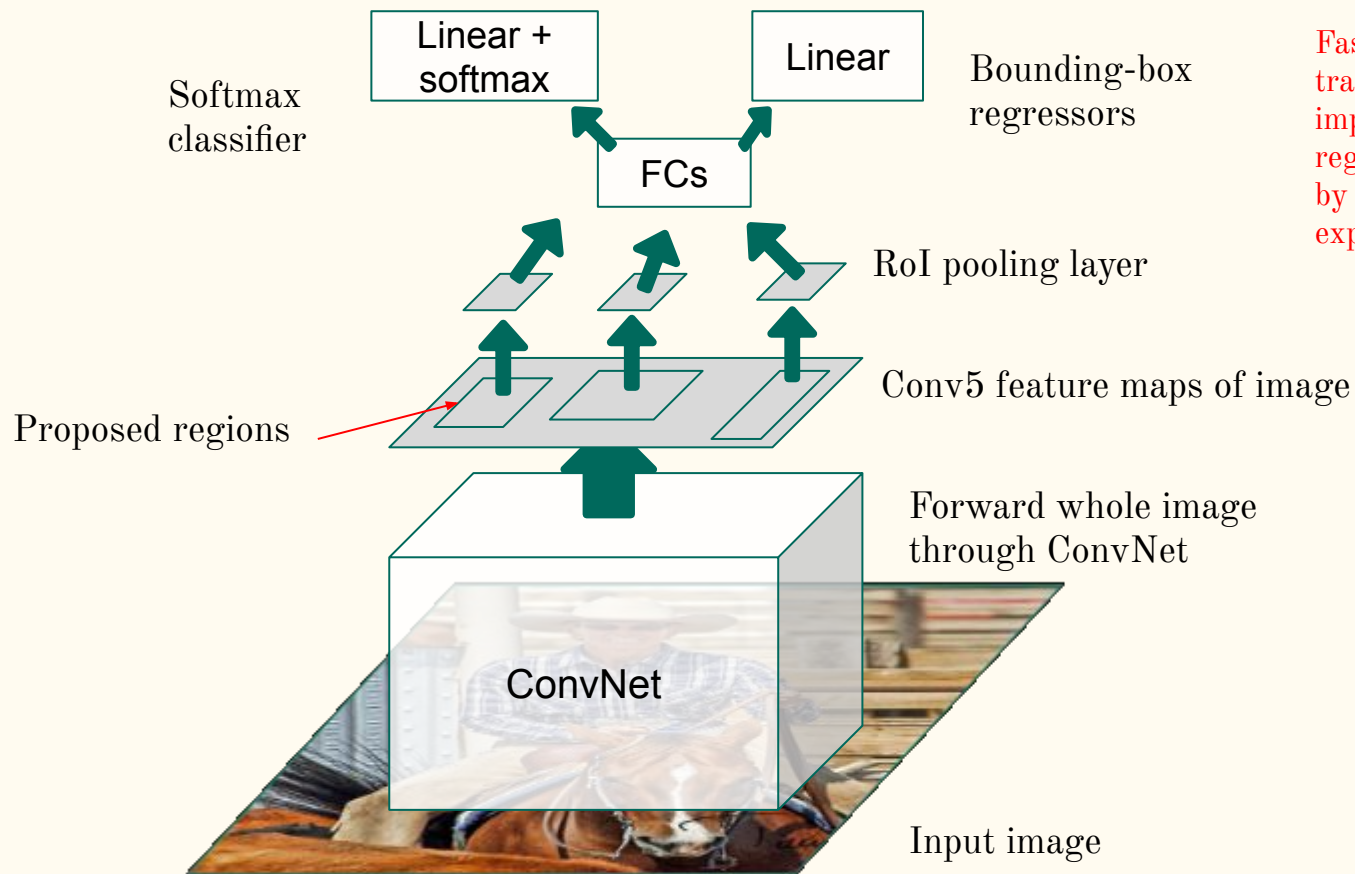
Fast R-CNN



Fast R-CNN



Fast R-CNN




Problem: Speed Bottleneck


Fast R-CNN is much faster in both training and testing time. However, the improvement is not dramatic because the region proposals are generated separately by another model and that is very expensive.

R-CNN VS SPP-Net VS Fast R-CNN

Training time (Hours)



Test time (seconds)




Faster R-CNN

Make CNN do proposals!

Insert Region Proposal Network (RPN) to predict proposals from features

Jointly train with 4 losses:

1. RPN classify object / not object
2. RPN regress box coordinates
3. Final classification score (object classes)
4. Final box coordinates




Faster R-CNN

Make CNN do proposals!


~ 7 FPS



Mask R-CNN


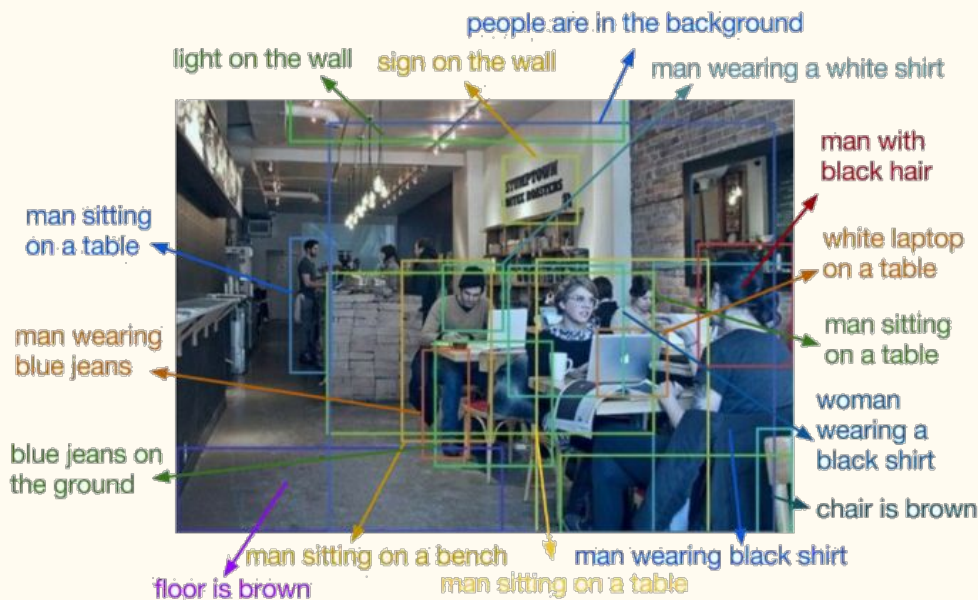


Mask R-CNN: Very good results

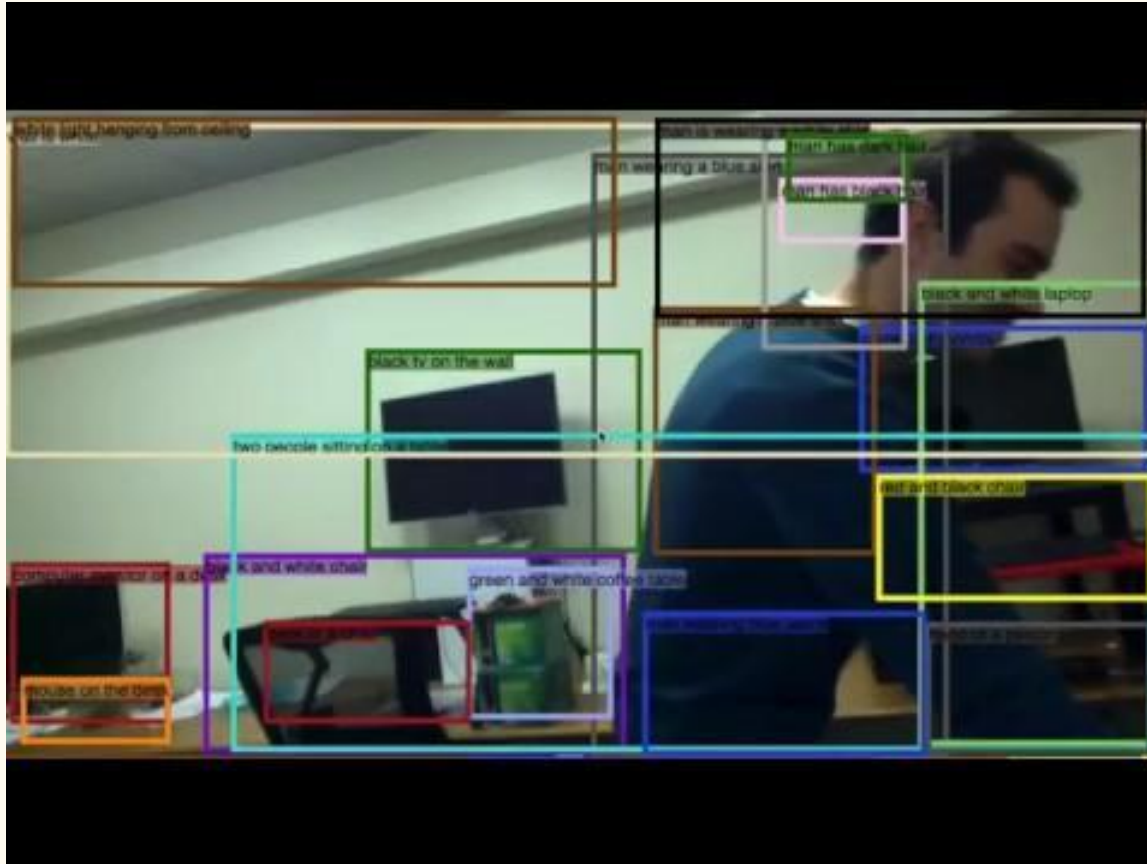


<https://arxiv.org/abs/1703.06870>

Aside: Object Detection + Captioning = Dense Captioning



Aside: Object Detection + Captioning = Dense Captioning



<https://arxiv.org/abs/1511.07571>

T.HANKS!

